Yash Sharma

syash2297@gmail.com | LinkedIn | GitHub | Portfolio | +91-9461572488

SUMMARY

AI/ML Engineer specializing in NLP, RAG pipelines, and Agentic AI systems. Proven track record of shipping end-to-end ML solutions with 93%+ F1-scores and sub-15ms inference latency using FastAPI, Docker, and CI/CD.

SKILLS

Languages: Python, SQL, JavaScript, TypeScript
ML & NLP: Scikit-learn, XGBoost, NumPy, Pandas, TF-IDF, LDA, K-Means, Feature Engineering, SHAP
LLM & GenAI: RAG, LangChain, LlamaIndex, Hugging Face Transformers, OpenAI API, pgvector, Hybrid Retrieval, Vertex AI, Gemini, Prompt Engineering, Agentic AI
MLOps & Tools: FastAPI, Docker, GitHub Actions, CI/CD, Streamlit, pytest, Git, Linux, REST APIs
Databases: PostgreSQL, SQLite, Firebase/Firestore, VectorDB

WORK EXPERIENCE

Celebal Technologies *May 2025 – Jul 2025*
Data Science Intern *Jaipur, Rajasthan*

- Engineered an unsupervised NLP pipeline (TF-IDF + K-Means + LDA) across 50,000+ records, reducing manual triage effort by 40%.
- Owned full ML lifecycle — data ingestion, feature engineering, model evaluation, and production deployment via FastAPI with CI/CD and Streamlit observability.

TaskUs *Feb 2022 – Sep 2023*
Teammate *Gurugram, Haryana*

- Delivered 95%+ CSAT scores and reduced repeat cases by 15% through data-driven issue pattern analysis.

Concentrix *Apr 2021 – Feb 2022*
Customer Service Representative *Gurugram, Haryana*

- Reduced repeat contacts by 15% via root-cause analysis on high-frequency issue patterns and structured resolution workflows.

PROJECTS

Customer Inquiry Classifier | Python, Scikit-learn, FastAPI, Docker, CI/CD, Streamlit

- Built an end-to-end NLP classifier (TF-IDF + LinearSVC) achieving 93%+ weighted F1-score across 7 classes; optimized inference to <15ms via FastAPI and containerized with Docker.
- Integrated GitHub Actions for automated testing and Streamlit dashboard for real-time performance monitoring and model observability.

TraceRAG — Full-Stack RAG System | FastAPI, Next.js, PostgreSQL, pgvector, LangChain

- Architected a hybrid retrieval system (dense vector + BM25 keyword) using pgvector, improving document retrieval precision over single-mode search.
- Engineered citation-grounded response generation with query tracing and faithfulness evaluation for hallucination reduction.

AI Trip Planner — Google Gen AI Exchange Hackathon | Gemini, Vertex AI, Node.js, React

- Participated in Google Gen AI Exchange Hackathon 2025; built real-time LLM streaming and multi-step prompt orchestration with Gemini on Vertex AI.

EDUCATION

Poddar Management and Technical Campus *Aug 2024 – Present*
Master of Computer Applications— Jaipur, Rajasthan

Singhania University *May 2020 – Jul 2023*
Bachelor of Computer Applications — Jhunjhunu, Rajasthan

CERTIFICATIONS

- **Google Gen AI Exchange Hackathon 2025** — Participant, Google (Gemini, Vertex AI, multi-step prompt orchestration)
- **Prime AI/ML Batch** — Apna College (Python, ML, DL, GenAI, AI Engineering) **[Pursuing]**
- **GenAI Mastermind** — Outskill (LLM workflows, AI agents, automation)